




Architecture — Training leverages parameter-efficient fine-tuning (PEFT) with LoRA to adapt the vision and language projection layers without modifying the massive base model. Bfloat16 precision prevents underflow issues common in visual token scaling.