




Execution — The vLLM engine hosts the merged model. Instead of stretching the image, Qwen's visual architecture dynamically tokens the image patches based on resolution, feeding them alongside the text prompt. The system instructions enforce a strict JSON structure for the backend to ingest.