




Architecture — The codebase uses a strict separation of concerns. Offline ML pipelines (data, train, eval) run independently from the online serving stack (gateway, vLLM). Configuration is injected globally via YAML.