




Training Dynamics — Transfer learning from COCO pre-trained weights significantly accelerates convergence. The engine tracks mAP@0.5:0.95 on the validation set at every epoch, automatically saving best.pt to prevent overfitting.