




Edge Optimization — Quantizing from 32-bit floats to 8-bit integers drastically reduces model size (usually by 4x) and speeds up inference on mobile CPUs/NPUs. The validation set is used during calibration to compute the scaling factors, minimizing the inevitable drop in accuracy.