



Qwen2.5-VL LoRA Fine-Tuning Dynamics

(a) Loss convergence


— Train loss — Eval loss



(b) Mean token accuracy




(c) Cosine learning-rate schedule



(d) Gradient norm stability

..... Clip threshold



Last logged train loss 0.0097; runtime 2.58 h