





(a) Loss convergence

— Train loss —● Eval loss

**(b) Mean token accuracy****(c) Cosine learning-rate schedule****(d) Gradient norm stability**

..... Clip threshold



Last logged train loss 0.0097; runtime 2.58 h