


VLM Training Pipeline Flow

PEFT / LoRA tuning of Qwen2.5-VL-3B with visual tokens, distributed via Accelerate, logging to W&B, culminating in adapter merging.



Architecture — Training leverages parameter-efficient fine-tuning (PEFT) with LoRA to adapt the vision and language projection layers without modifying the massive base model. Bfloat16 precision prevents underflow issues common in visual token scaling.