


VLM Inference Pipeline Flow

Local generation with vLLM server: Image parsing, dynamic visual tokenization, LLM text generation, and JSON constraint extraction.



Execution — The vLLM engine hosts the merged model. Instead of stretching the image, Qwen's visual architecture dynamically tokens the image patches based on resolution, feeding them alongside the text prompt. The system instructions enforce a strict JSON structure for the backend to ingest.