


VLM System Architecture Overview

The modular VLM architecture separates offline data and training workflows from the online gateway and inference-serving stack.



Architecture — The codebase uses a strict separation of concerns. Offline ML pipelines (data, train, eval) run independently from the online serving stack (gateway, vLLM). Configuration is injected globally via YAML.