


YOLO Deployment Architecture Flow

Dual deployment strategy: on-device edge inference for real-time overlays, with an optional server fallback for heavier model execution and centralized logging.



Strategy — The primary detection runs locally on the mobile device (Edge) via TFLite, providing instant visual feedback without network latency. The backend ONNX Runtime serves as an optional high-accuracy fallback or for logging diagnostic imagery when device constraints prevent local processing.