


YOLO System Architecture Overview

The high-level architecture of the YOLO computer vision pipeline, separating offline training mechanics from the exported edge artifacts.



Architecture — The YOLO architecture is cleanly split. The ML lifecycle handles the heavy lifting of data preparation and model training using PyTorch. The resultant artifacts are quantized and exported, decoupling the Python-based training from the lightweight C++/Java/Swift runtimes used on mobile edges.